

Problem Set 3

web.stanford.edu/class/stats202/content/viewhw.html?hw3

Problem 1

Chapter 4, Exercise 4 (Sec. 4.7, p. 168)

Part A

- x is uniformly distributed on $[0, 1]$.
- When predicting a test observation's response, we look the 10% of the range closest to that observation.
 - If our test observation has value $x = 0.6$, we look at $[0.55, 0.65]$.
 - If our test observation has value $x = 0.02$, we look at $[0.00, 0.10]$.
 - If our test observation has value $x = 0.98$, we look at $[0.90, 1.00]$.

Since at any give point we're looking at 10% of the range and the points are evenly distributed along that range, we'd expect to be looking at 10% of the data on average each time.

Part B

- x_1 is uniformly distributed on $[0, 1]$, and x_2 is uniformly distributed on $[0, 1]$.
- Similar rules as in part a.

Since at any give point we're looking at 10% of x_1 's range and 10% of x_2 's range and the points are evenly distributed along those two ranges, we'd expect to be looking at 1% of the data on average each time.

We can think about it as a square:

	0	1	2	3	4	5	6	7	8	9
0	-		-		-		-		-	
1	-		-		-		-		-	
2	-		-		-		-		-	
3	-		-		-		X		-	
4	-		-		-		-		-	
5	-		-		-		-		-	
6	-		-		-		-		-	

7	-	-	-	-	-	-	-	-	-	-	-	-
8	-	-	-	-	-	-	-	-	-	-	-	-
9	-	-	-	-	-	-	-	-	-	-	-	-

If we look at just 10% of the x-dimension (let's say the 5th column) and then also just 10% of the y-dimension (let's say the 3rd row), we get $1/100 = 1\%$ of the available cells.

Part C

If we have 100 features (a.k.a. 100 dimensions) and we look at just 10% of the range for each of them, we look at just a tiny portion (10^{-100}) of the data.

Part D

Let's say we have 1 billion (10^9) training observations. That's a lot of data! However, consider trying to predict the response for some test observation m with 100 features, where we look at just the observations that fall within 10% of each range from m . Of the 1 billion points we started out with, we'd expect to have $10^9 \cdot 10^{-100} = 10^{-91}$ observations to look at. That is still effectively 0, which doesn't help us at all.

Part E

The expected length of the hypercube is:

- hypercube's length is $(\frac{1}{10})^1 = 10\%$ when $p = 1$
- hypercube's length is $(\frac{1}{10})^2 = 1\%$ when $p = 2$
- hypercube's length is $(\frac{1}{10})^{100}$ when $p = 100$

Problem 2

Chapter 4, Exercise 6 (Sec. 4.7, p. 170).

$$p(X) = \frac{e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}}.$$

Part A

X1 = hours studied, X2 = undergrad GPA, and Y = receive an A. We fit a logistic regression and produce estimated coefficient, $\beta^0 = -6$, $\beta^1 = 0.05$, $\beta^2 = 1$.

$$\begin{aligned}
 x &= \text{"3.5 GPA \& studies for 40h"} \\
 Pr(x) &= 0.5 = \frac{e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2}}{1 + e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2}} \\
 &= \frac{e^{-6 + 40 \cdot 0.05 + 3.5 \cdot 1.0}}{1 + e^{-6 + 40 \cdot 0.05 + 3.5 \cdot 1.0}} \\
 &= \boxed{0.377541}
 \end{aligned}$$

Part B

$$\begin{aligned}
 x &= \text{"3.5 GPA \& studies for 40h"} \\
 Pr(x) &= 0.5 = \frac{e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2}}{1 + e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2}} \\
 &= \frac{e^{-6 + 0.05 \cdot \text{num_hours} + 3.5 \cdot 1.0}}{1 + e^{-6 + 0.05 \cdot \text{num_hours} + 3.5 \cdot 1.0}} \\
 &= \boxed{50 \text{ hours}}
 \end{aligned}$$

Input interpretation:

solve	$0.5 = \frac{e^{-6 + 0.05x + 3.5}}{1 + e^{-6 + 0.05x + 3.5}}$	for	x
-------	---	-----	---

Result:

$$x = \underline{50 + (125.664 i) n} \text{ and } n \in \mathbb{Z}$$

Problem 3

Chapter 4, Exercise 8 (Sec. 4.7, p. 170).

We prefer to use the logistic regression, despite the fact that the 1-nearest neighbors method gives us a lower average error rate than the logistic regression (18% vs $\frac{20+30}{2} = 25\%$). However, since KNN with $k = 1$ gives us a training error rate of 0% , we know that its tests error rate must be 36% (since $\frac{0+x}{2} = 18$).

Problem 4

Chapter 4, Exercise 10 (Sec. 4.7, p. 171). In part (i), please be concise; only describe and provide the output of your best prediction.

```
##           Year      Lag1      Lag2      Lag3      Lag4
## Year      1.00000000 -0.032289274 -0.03339001 -0.03000649 -0.031127923
## Lag1     -0.03228927  1.000000000 -0.07485305  0.05863568 -0.071273876
## Lag2     -0.03339001 -0.074853051  1.00000000 -0.07572091  0.058381535
## Lag3     -0.03000649  0.058635682 -0.07572091  1.00000000 -0.075395865
## Lag4     -0.03112792 -0.071273876  0.05838153 -0.07539587  1.000000000
## Lag5     -0.03051910 -0.008183096 -0.07249948  0.06065717 -0.075675027
## Volume    0.84194162 -0.064951313 -0.08551314 -0.06928771 -0.061074617
## Today    -0.03245989 -0.075031842  0.05916672 -0.07124364 -0.007825873
##           Lag5      Volume      Today
## Year    -0.030519101  0.84194162 -0.032459894
## Lag1    -0.008183096 -0.06495131 -0.075031842
## Lag2    -0.072499482 -0.08551314  0.059166717
## Lag3     0.060657175 -0.06928771 -0.071243639
## Lag4    -0.075675027 -0.06107462 -0.007825873
## Lag5     1.000000000 -0.05851741  0.011012698
## Volume  -0.058517414  1.00000000 -0.033077783
## Today    0.011012698 -0.03307778  1.000000000
```

Part B

Only the `Lag2` predictor appears statistically significant.

```
## (Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
## 0.001898848 0.118144368 0.029601361 0.546923890 0.293653342 0.583348244
##           Volume
## 0.537674762
```

Part C

```
attach(Weekly)
probs = predict(fit, type = 'response')
pred = rep('Down', nrow(Weekly))
pred[probs > .5] = 'Up'

table(pred, Direction)
```

```
##           Direction
## pred   Down  Up
##   Down    54  48
##    Up   430 557
```

```
mean(pred==Direction) # Fraction of correct predictions
```

```
## [1] 0.5610652
```

This tells us that we are making the correct prediction about 56% of the time. In particular, we often wrongly predict “Up” when we should have predicted “Down”.

Part D

```
train=(Year<=2008)
Weekly.2009and10 = Weekly[!train,]
dim(Weekly.2009and10) # 0 9
```

```
## [1] 104 9
```

```
Direction.2009and10 = Direction[!train]

fit = glm(Direction ~ Lag2, data = Weekly, family = binomial, subset = train)
probs = predict(fit, Weekly.2009and10, type = 'response')

pred = rep('Down', nrow(Weekly.2009and10))
pred[probs > .5] = 'Up'
table(pred, Direction.2009and10)
```

```
##           Direction.2009and10
## pred   Down Up
##   Down    9  5
##    Up   34 56
```

```
mean(pred==Direction.2009and10)
```

```
## [1] 0.625
```

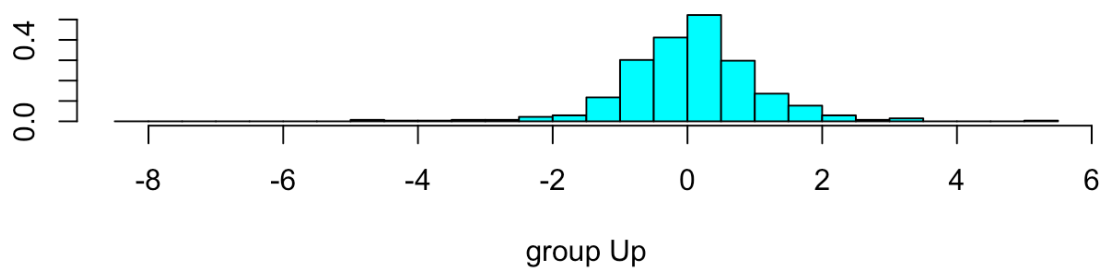
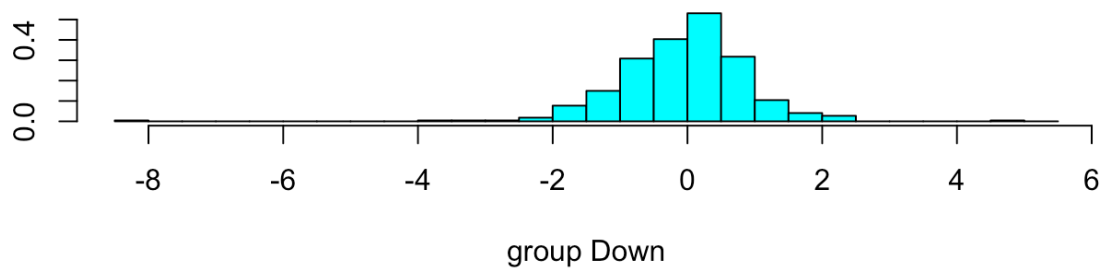
Using just `Lag2` gives us a better result of `62.5%`.

Part E

```
library(MASS)
lda.fit = lda(Direction ~ Lag2, data = Weekly, subset = train)
lda.fit
```

```
## Call:
## lda(Direction ~ Lag2, data = Weekly, subset = train)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag2
## Down -0.03568254
## Up    0.26036581
##
## Coefficients of linear discriminants:
##      LD1
## Lag2 0.4414162
```

```
plot(lda.fit)
```



```
lda.pred = predict(lda.fit, Weekly.2009and10)
names(lda.pred)
```

```
## [1] "class"      "posterior" "x"
```

```
lda.class = lda.pred$class
table(lda.class, Direction.2009and10)
```

```
##           Direction.2009and10
## lda.class Down Up
##      Down      9  5
##      Up      34 56
```

```
mean(lda.class == Direction.2009and10)
```

```
## [1] 0.625
```

```
sum(lda.pred$posterior[,1] >= .5)
```

```
## [1] 14
```

```
sum(lda.pred$posterior[,1] < .5)
```

```
## [1] 90
```

Part F

```
qda.fit = qda(Direction ~ Lag2, data = Weekly, subset = train)
qda.fit
```

```
## Call:
## qda(Direction ~ Lag2, data = Weekly, subset = train)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag2
## Down -0.03568254
## Up    0.26036581
```

```
qda.class = predict(qda.fit, Weekly.2009and10)$class
table(qda.class, Direction.2009and10)
```



```
##           Direction.2009and10
## qda.class Down Up
##      Down   0  0
##      Up    43 61
```

```
mean(qda.class == Direction.2009and10)
```

```
## [1] 0.5865385
```

Part G

```
library(class)
train.X = as.matrix(Lag2[train])
test.X  = as.matrix(Lag2[!train])
train.Direction = Direction[train]

set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 1)
table(knn.pred, Direction.2009and10)
```

```
##           Direction.2009and10
## knn.pred Down Up
##      Down   21 30
##      Up    22 31
```

```
mean(knn.pred == Direction.2009and10)
```

```
## [1] 0.5
```

Part H

Logistic Regression and Linear Discriminant Analysis tied for the best test results, both resulting in a 62.5% success rate.

Part I

2-means

```
train.X = as.matrix(Lag2[train])
test.X  = as.matrix(Lag2[!train])
train.Direction = Direction[train]

set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 2)
table(knn.pred, Direction.2009and10)
```

```
##           Direction.2009and10
## knn.pred Down Up
##      Down   19 27
##      Up    24 34
```

```
mean(knn.pred == Direction.2009and10)
```

```
## [1] 0.5096154
```

3-means

```
train.X = as.matrix(Lag2[train])
test.X  = as.matrix(Lag2[!train])
train.Direction = Direction[train]

set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 3)
table(knn.pred, Direction.2009and10)
```

```
##           Direction.2009and10
## knn.pred Down Up
##      Down   16 20
##      Up    27 41
```

```
mean(knn.pred == Direction.2009and10)
```

```
## [1] 0.5480769
```

4-means

```
train.X = as.matrix(Lag2[train])
test.X  = as.matrix(Lag2[!train])
train.Direction = Direction[train]

set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 4)
table(knn.pred, Direction.2009and10)
```

```
##           Direction.2009and10
## knn.pred Down Up
##      Down   20 17
##      Up    23 44
```

```
mean(knn.pred == Direction.2009and10)
```

```
## [1] 0.6153846
```

5-means

```
train.X = as.matrix(Lag2[train])
test.X = as.matrix(Lag2[!train])
train.Direction = Direction[train]

set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 5)
table(knn.pred, Direction.2009and10)
```

```
##           Direction.2009and10
## knn.pred Down Up
##      Down   16 21
##      Up    27 40
```

```
mean(knn.pred == Direction.2009and10)
```

```
## [1] 0.5384615
```

Problem 5

```
library(MASS)
library(nnet)
library(ggplot2)

rosters <- read.csv('rosters.csv')
summary(rosters)
```

```
##           X           gender           height           homestate
## Min.      : 0.00   male:204   Min.      :62.00   CA       :70
## 1st Qu.: 50.75                1st Qu.:71.00   WA       :12
## Median :101.50                Median :74.00   TX       :11
## Mean    :101.63                Mean    :73.33   GA       : 9
## 3rd Qu.:152.25                3rd Qu.:76.00   AZ       : 8
## Max.    :206.00                Max.    :84.00   FL       : 8
##                                     (Other):86
##           name           sport           weight
## Alabi, Adrian      : 1   Baseball   :29   Min.    :125.0
## Alexander, Terrence: 1   Basketball:12   1st Qu.:176.5
## Alfieri, Joey       : 1   Football  :95   Median  :197.0
## Allen, Malcolm      : 1   Soccer    :25   Mean    :205.4
## Allen, Marcus       : 1   Tennis    :10   3rd Qu.:229.2
## Allen, Rosco        : 1   Wrestling :33   Max.    :321.0
## (Other)             :198
```

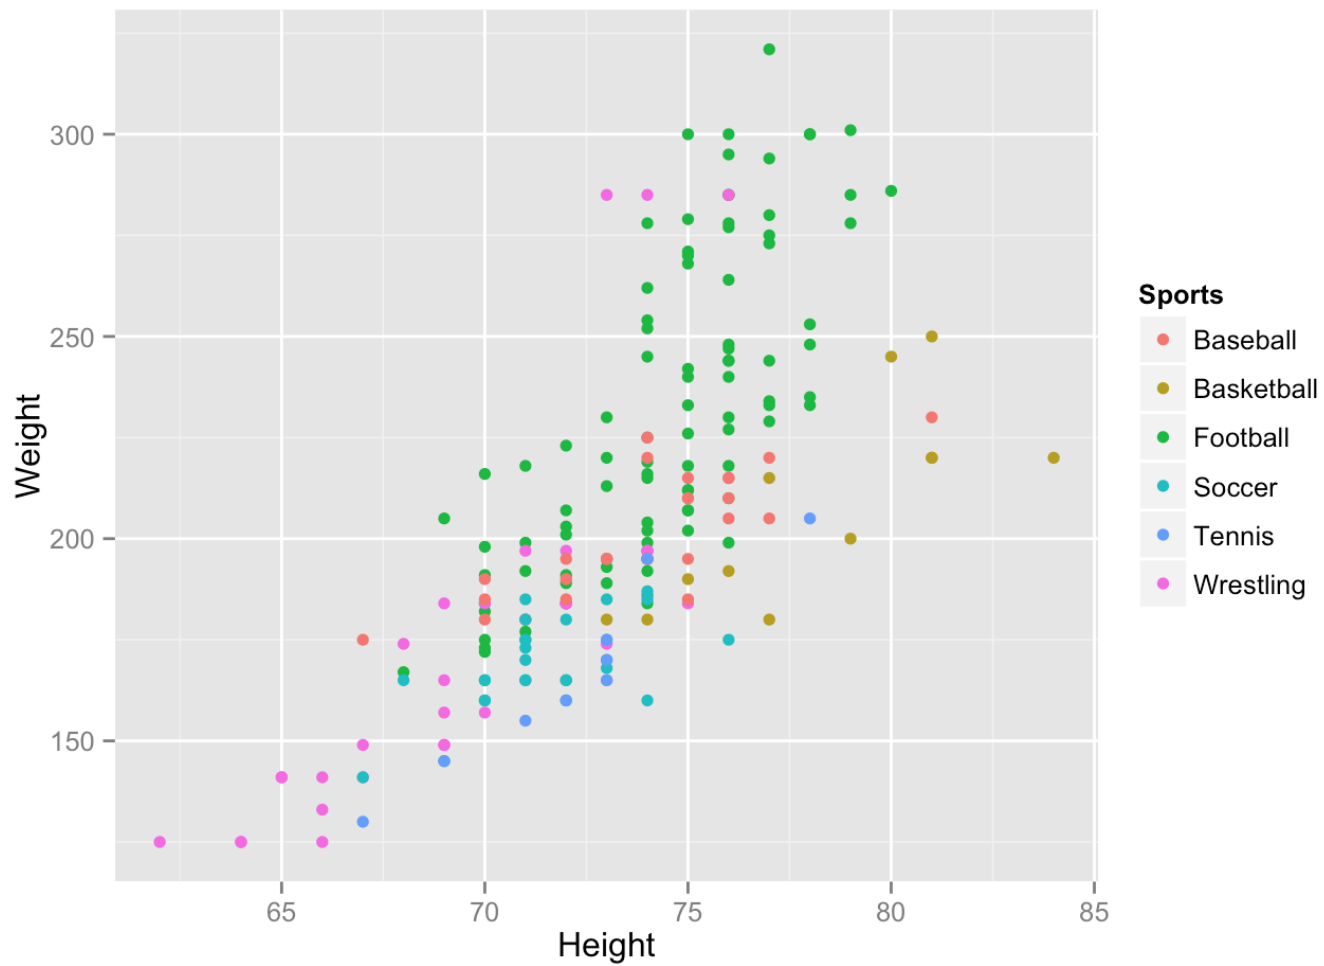
```
fit = multinom(sport ~ height + weight, rosters)
```

```
## # weights:  24 (15 variable)
## initial  value 365.518932
## iter   10 value 266.164885
## iter   20 value 207.817029
## iter   30 value 200.902009
## iter   40 value 200.198801
## iter   50 value 199.808004
## iter   60 value 199.680438
## iter   70 value 199.653822
## iter   80 value 199.647074
## iter   90 value 199.645164
## final   value 199.644727
## converged
```

```
summary(fit)
```

```
## Call:
## multinom(formula = sport ~ height + weight, data = rosters)
##
## Coefficients:
##           (Intercept)      height      weight
## Basketball  -71.754980   1.1940557 -0.09567987
## Football     16.366376  -0.3500011  0.05057577
## Soccer       -6.939543   0.3534022 -0.10321321
## Tennis      -30.636220   0.8123257 -0.16392309
## Wrestling    44.531678  -0.6709520  0.02033305
##
## Std. Errors:
##           (Intercept)      height      weight
## Basketball  0.140277235  0.06577436  0.02461660
## Football    4.953425402  0.08643443  0.01180643
## Soccer      0.083352009  0.04979394  0.01992204
## Tennis      0.004621523  0.07325690  0.03054070
## Wrestling   3.278747021  0.06415213  0.01411842
##
## Residual Deviance: 399.2895
## AIC: 429.2895
```

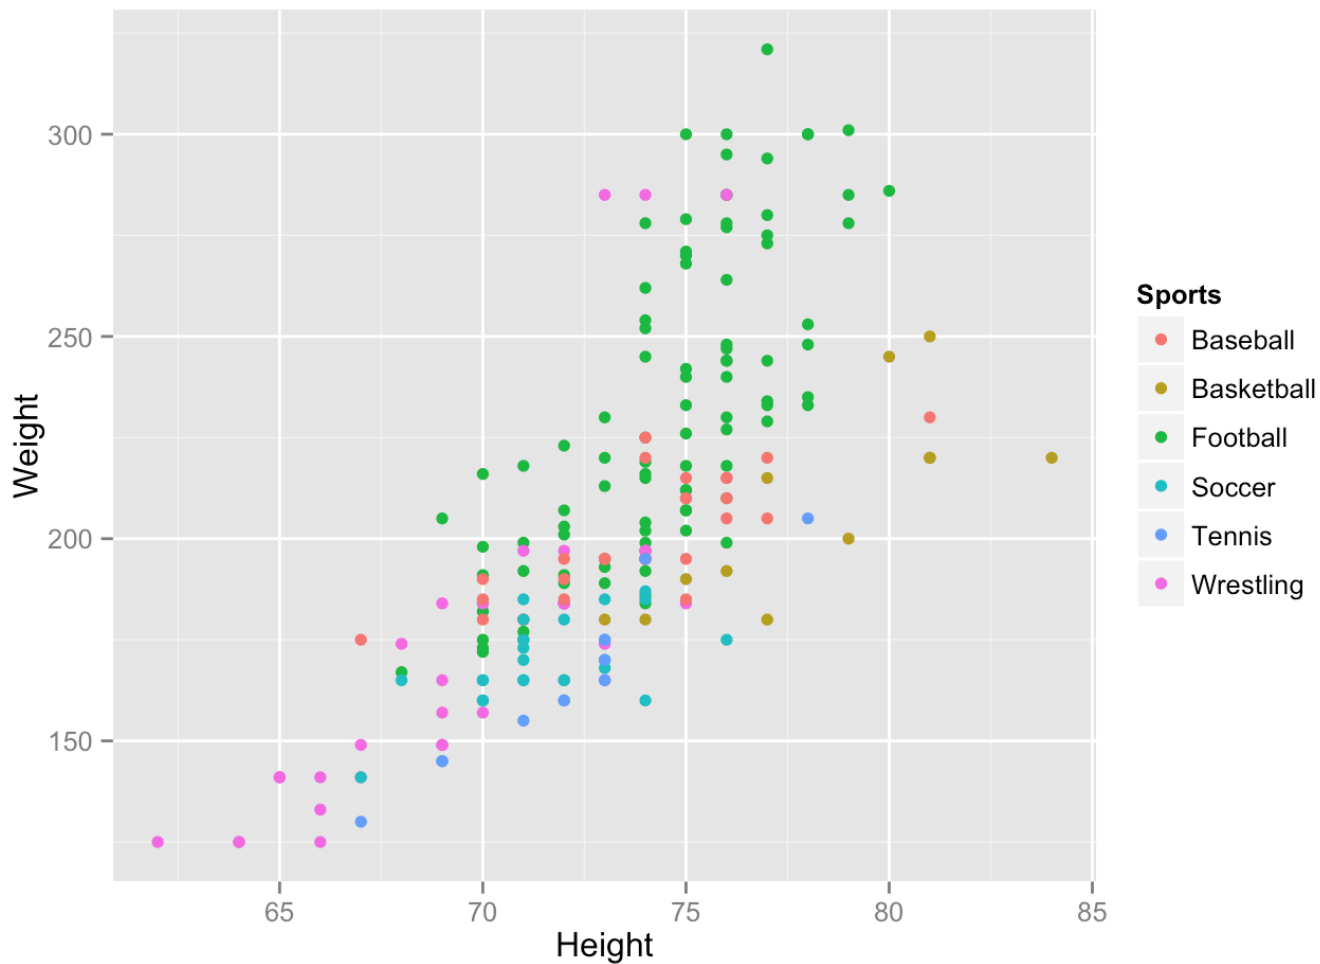
```
Height = rosters$height
Weight = rosters$weight
Sports = rosters$sport
qplot(Height, Weight, rosters, color=Sports)
```



```
pred = predict(fit, rosters)
pred
```

##	[1]	Wrestling	Football	Football	Football	Football
##	[6]	Football	Soccer	Football	Football	Football
##	[11]	Football	Wrestling	Football	Football	Football
##	[16]	Football	Football	Football	Football	Football
##	[21]	Football	Football	Football	Baseball	Football
##	[26]	Football	Football	Football	Football	Football
##	[31]	Football	Wrestling	Football	Football	Football
##	[36]	Wrestling	Football	Football	Soccer	Football
##	[41]	Football	Football	Football	Football	Football
##	[46]	Football	Football	Wrestling	Football	Football
##	[51]	Wrestling	Football	Football	Football	Football
##	[56]	Football	Football	Football	Football	Football
##	[61]	Football	Football	Football	Football	Football
##	[66]	Football	Football	Football	Football	Football
##	[71]	Football	Football	Football	Football	Football
##	[76]	Football	Football	Football	Football	Football
##	[81]	Football	Football	Wrestling	Football	Football
##	[86]	Wrestling	Football	Football	Football	Football
##	[91]	Football	Football	Football	Football	Football
##	[96]	Football	Soccer	Wrestling	Football	Soccer
##	[101]	Football	Football	Wrestling	Wrestling	Wrestling
##	[106]	Wrestling	Soccer	Football	Football	Wrestling
##	[111]	Soccer	Football	Football	Wrestling	Soccer
##	[116]	Soccer	Football	Wrestling	Wrestling	Football
##	[121]	Soccer	Wrestling	Wrestling	Football	Wrestling
##	[126]	Football	Wrestling	Wrestling	Basketball	Soccer
##	[131]	Soccer	Basketball	Football	Basketball	Baseball
##	[136]	Basketball	Basketball	Football	Basketball	Basketball
##	[141]	Football	Wrestling	Football	Football	Wrestling
##	[146]	Football	Football	Football	Football	Wrestling
##	[151]	Football	Football	Football	Football	Football
##	[156]	Football	Football	Football	Soccer	Football
##	[161]	Football	Football	Baseball	Football	Basketball
##	[166]	Football	Basketball	Football	Football	Football
##	[171]	Wrestling	Wrestling	Soccer	Soccer	Soccer
##	[176]	Wrestling	Football	Football	Soccer	Soccer
##	[181]	Soccer	Football	Soccer	Soccer	Tennis
##	[186]	Soccer	Soccer	Wrestling	Wrestling	Tennis
##	[191]	Soccer	Wrestling	Baseball	Soccer	Basketball
##	[196]	Soccer	Soccer	Soccer	Football	Soccer
##	[201]	Soccer	Soccer	Soccer	Soccer	
##	Levels:	Baseball	Basketball	Football	Soccer	Tennis Wrestling

```
qplot(Height, Weight, pred, color=Sports)
```

```
table(rosters$sport, pred)
```

```
##          pred
##      Baseball Basketball  Football  Soccer  Tennis  Wrestling
## Baseball          1           2        22         1         0          3
## Basketball         1           7         2         2         0          0
## Football           1           0        84         2         0          8
## Soccer             1           0         4        12         2          6
## Tennis             0           1         1         8         0          0
## Wrestling          0           0        12         7         0         14
```

```
success_rate = mean(pred == rosters$sport)
```

```
## Our success rate is 0.5784314
```

```
## Our 0-1 loss error rate is 0.4215686
```