

Problem 3

Chapter 5, Exercise 8 (Sec. 5.4, p. 200)


Part A

```
set.seed(1)
y = rnorm(100)
x = rnorm(100)
y = x - 2*x^2 + rnorm(100)
```

```
n = 100  # number of samples
p = 2    # number of dimensions
```

Part B

```
qplot(x, y)
```



There is clearly a non-random relationship between x and y . In particular, y takes on a parabolic shape centered around 0 as x varies. However, there is a range of variance of a bit over 1, creating a band of values rather than a neat line of points.

Part C

```
data = data.frame(y = y, x = x)
```

Linear:

```
model = glm(y ~ x, data = data)
model$coef
```

```
## (Intercept)          x
## -1.8185184    0.2430443
```

```
## standard k-fold CV estimate = 5.890979
##      bias-corrected version = 5.888812
```

Squared:

```
model = glm(y ~ poly(x, degree = 2), data = data)
model$coef
```

```
##      (Intercept) poly(x, degree = 2)1 poly(x, degree = 2)2
##      -1.827707      2.316401      -21.058587
```

```
## standard k-fold CV estimate = 1.086596
##      bias-corrected version = 1.086326
```

Cubic:

```
model = glm(y ~ poly(x, degree = 3), data = data)
model$coef
```

```
##      (Intercept) poly(x, degree = 3)1 poly(x, degree = 3)2
##      -1.8277074      2.3164010      -21.0585869
## poly(x, degree = 3)3
##      -0.3048398
```

```
## standard k-fold CV estimate = 1.102585
##      bias-corrected version = 1.102227
```

Quadratic:

```
model = glm(y ~ poly(x, degree = 4), data = data)
model$coef
```

```
##          (Intercept) poly(x, degree = 4)1 poly(x, degree = 4)2
##          -1.8277074          2.3164010          -21.0585869
## poly(x, degree = 4)3 poly(x, degree = 4)4
##          -0.3048398          -0.4926249
```

```
## standard k-fold CV estimate = 1.114772
##          bias-corrected version = 1.114334
```

Part D

```
set.seed(5)
y = rnorm(100)
x = rnorm(100)
y = x - 2*x^2 + rnorm(100)

data = data.frame(y = y, x = x)
```

Linear:

```
model = glm(y ~ x, data = data)
```

```
## standard k-fold CV estimate = 10.32995
##          bias-corrected version = 10.32529
```

Squared:

```
model = glm(y ~ poly(x, degree = 2), data = data)
```

```
## standard k-fold CV estimate = 0.9586209
##          bias-corrected version = 0.9583222
```

Cubic:

```
model = glm(y ~ poly(x, degree = 3), data = data)
```

```
## standard k-fold CV estimate = 0.9867481
##          bias-corrected version = 0.9862594
```

Quadratic:

```
model = glm(y ~ poly(x, degree = 4), data = data)
```

```
## standard k-fold CV estimate = 1.335795
##          bias-corrected version = 1.332331
```

Both runs use the same data generator function, but the `rnorm(...)` creates variance between runs. This variance results in slightly different fits resulting from a linear regression model.

Part E

The squared model had the smallest error. This is what I expected, since the original function is based off the square of x .

Part F

Our original function was $y = x - 2x^2 + \text{rnorm}(100)$, which means the correct coefficients ought to be: $B_0 = 0$, $B_1 = 1$, $B_2 = -2$.

Instead, we got $[-1.82, 0.24]$, $[-1.83, 2.32, -21.06]$, $[-1.83, 2.31, -21.06, -0.35]$, and $[-1.83, 2.32, -21.06, -0.31, -0.49]$ for the linear, squared, cubed, and quadratic fits respectively. These are wayyyy off, even for the best fit (squared, with coefficients $[-1.83, 2.32, -21.06]$). This does not agree with the conclusions drawn based on the cross-validation results, which implied that the squared fit was fairly accurate.